Database Constraints and Design

- We know that databases are often required to satisfy some *integrity constraints.*

- The most common ones are functional and inclusion dependencies.

- We'll study properties of integrity constraints.

- We show how they influence database design, in particular, how they tell us what type of information should be recorded in each particular relation

343   1   Introduction to Databases

Review: functional dependencies and keys

- A functional dependency is $X \rightarrow Y$ where X, Y are sequences of attributes. It holds in a relation R if for every two tuples t_1, t_2 in R:

$$\pi_X(t_1) = \pi_X(t_2) \quad \text{implies} \quad \pi_Y(t_1) = \pi_Y(t_2)$$

- A very important special case: *keys*

- Let K be a set of attributes of R, and U the set of *all* attributes of R. Then K is a key if R satisfies functional dependency $K \rightarrow U$.

- In other words, a set of attributes K is a key in R if for any two tuples t_1, t_2 in R,

$$\pi_K(t_1) = \pi_K(t_2) \quad \text{implies} \quad t_1 = t_2$$

- That is, a key is a set of attributes that uniquely identify a tuple in a relation.

343   2   Introduction to Databases

Problems

- Good constraints vs Bad constraints: some constraints may be undesirable as they lead to problems.

- Implication problem: Suppose we are given some constraints. Do they imply others?

  This is important, as we never get the list of all constraints that hold in a database (such a list could be very large, or just unknown to the designer).

  It is possible that all constraints given to us look OK, but they imply some bad constraints.

- Axiomatization of constraints: a simple way to state when some constraints imply others.

Bad constraints: example

DME	Dept	Manager	Employee
	D1	Smith	Jones
	D1	Smith	Brown
	D2	Turner	White
	D3	Smith	Taylor
	D4	Smith	Clarke
	...	...	...

ES	Employee	Salary
	Jones	10
	Brown	20
	White	20
	...	...

- Functional dependencies

- in DME: Dept $\rightarrow$ Manager, but none of the following:

  Dept $\rightarrow$ Employee

  Manager $\rightarrow$ Dept

  Manager $\rightarrow$ Employee

- in ES: Employee $\rightarrow$ Salary (Employee is a key for ES)

  but not Salary $\rightarrow$ Employee

Update anomalies

- **Insertion anomaly**: A company hires a new employee, but doesn't immediately assign him/her to a department. We cannot record this fact in relation DME.

- **Deletion anomaly**: Employee White leaves the company. We have to delete a tuple from DME. But this results in deleting manager Turner as well, even though he hasn't left!

- Reason for anomalies: the association between managers and employees is represented in the same relation as the association between managers and departments. Furthermore, the same fact (a department D is managed by M) could be represented more than once.

- Putting this into the language of functional dependencies: we have a dependency Dept $\rightarrow$ Manager, but Dept is not a key.

- In general, one tries to avoid situations like this.

dependencies and implication

- Suppose we have relation R with 3 attributes A, B, C

- Someone tells us that R satisfies $A \rightarrow B$ and $B \rightarrow \{A, C\}$.

- This looks like the bad case before: we have $A \rightarrow B$ that does not mention C.

- But A is a key: if $\pi_A(t_1) = \pi_A(t_2)$, then $\pi_B(t_1) = \pi_B(t_2)$, and since B is a key, $t_1 = t_2$.

- Thus, even though we did not have information that A is a key, we were able to derive it, as it is implied by other dependencies.

- Using implication, we can find all dependencies and figure out if there is a problem with the design.

Implication for functional dependencies

- Normally we try to find nontrivial dependencies

    A trivial dependency is $X \to Y$ with $Y \subseteq X$

- Suppose we are given a set U of attributes of a relation, a set F of functional dependencies (FDs), and a FD f over U.

- F implies f, written
$$F \vdash f$$
if for every relation over attributes U, if R satisfies all FDs in F, then it is the case that it also satisfies f.

- Problem: Given U and F, find all nontrivial FDs f such that $F \vdash f$.

Implication for functional dependencies cont'd

- Closed set with respect to F: a subset V of U such that, for every $X \to Y$ in F,
$$\text{if } X \subseteq V, \text{ then } Y \subseteq V$$

- Property of FDs: For every set V, there exists a unique set $C_F(V)$, called the *closure* of V with respect to F, such that

    $C_F(V)$ is closed

    $V \subseteq C_F(V)$

    For every closed set W, $V \subseteq W$ implies $C_F(V) \subseteq W$.

- Solution to the implication problem:

    A FD $X \to Y$ is implied by F
    if and only if
    $Y \subseteq C_F(X)$

Implication and closure

- To solve the implication problem, it suffices to find closure of each set of attributes.

- A naive approach to finding $C_F(X)$: check all subsets $V \subseteq U$, verify is they are closed, and select the smallest closed set containing X.

- Problem: this is too expensive.

- If U has n elements and X has $m < n$ elements, then there are 2^{n-m} subsets of U that contain V.

- But there is a very fast, $O(n)$, algorithm to compute $C_F(X)$.

- We will see a very simple $O(n^2)$ algorithm instead.

Implication and closure cont'd

ALGORITHM CLOSURE

Input: a set F of FDs, and a set X
Output: $C_F(X)$

1. $unused := F$
2. $closure := X$
3. repeat until no change:
    if $Y \rightarrow Z \in unused$ and $Y \subseteq closure$ then
        (i) $unused := unused - \{Y \rightarrow Z\}$
        (ii) $closure := closure \cup Z$
4. Output $closure$.

Homework: Prove that CLOSURE returns $C_F(X)$
and that its running time is $O(n^2)$.

Properties of the closure

- $X \subseteq C_F(X)$

- $X \subseteq Y$ implies $C_F(X) \subseteq C_F(Y)$

- $C_F(C_F(X)) = C_F(X)$

- It turns out that for any operator C satisfying

$$X \subseteq C(X)$$
$$X \subseteq Y \text{ implies } C(X) \subseteq C(Y)$$
$$C(C(X)) = C(X)$$

  there exists a family of FDs F such that $C = C_F$.

Closure: Example

- Common practice: write sets as AB for $\{A, B\}$, BC for $\{B, C\}$ etc

- $U = \{A, B, C\}$, $F = \{A \rightarrow B, B \rightarrow AC\}$

- Closure:

$$C_F(\emptyset) = \emptyset$$
$$C_F(C) = C$$
$$C_F(B) = ABC$$
hence $C_F(X) = ABC$ for any X that contains B
$$C_F(A) = ABC$$
hence $C_F(X) = ABC$ for any X that contains A

Keys, candidate keys, and prime attributes

- A set X of attributes is a key with respect to F if $X \rightarrow U$ is implied by F

- That is, X is a key if $C_F(X) = U$

- In the previous example: any set containing A or B is a key.

- *Candidate keys*: smallest keys. That is, keys X such that for each $Y \subset X$, Y is not a key.

- In the previous example: A and B.

- Suppose U has n attributes. What is the maximum number of candidate keys?

  answer: $\binom{n}{\lfloor n/2 \rfloor}$

  and that's very large: 252 for $n = 10$, $> 180,000$ for $n = 20$

- *Prime attribute*: an attribute of a candidate key.

Inclusion dependencies

- Functional and inclusion dependencies are the most common ones encountered in applications

- Reminder: $R[A_1, \ldots, A_n] \subseteq S[B_1, \ldots, B_n]$ says that for any tuple t in R, there is a tuple t' in S such that

$$\pi_{A_1,\ldots,A_n}(t) \quad = \pi_{B_1,\ldots,B_n}(t')$$

- Suppose we have a set of relation names and inclusion dependencies (IDs) between them. Can we derive *all* IDs that are valid?

- Formally, let G be a set of IDs and g and ID. We say that G implies g, written

$$G \vdash g$$

if any set of relations that satisfies all IDs in G, also satisfied g.

- The answer is positive, just as for FDs.

Implication of inclusion dependencies

- There are simple rules:

- $R[A] \subseteq R[A]$

-
$$\text{if} \quad R[A_1, \ldots, A_n] \subseteq S[B_1, \ldots, B_n]$$
$$\text{then} \quad R[A_{i_1}, \ldots, A_{i_n}] \subseteq S[B_{i_1}, \ldots, B_{i_n}]$$
where $\{i_1, \ldots, i_n\}$ is a permutation of $\{1, \ldots, n\}$.

-
$$\text{if } R[A_1, \ldots, A_m, A_{m+1}, \ldots, A_n] \subseteq S[B_1, \ldots, B_m, B_{m+1}, \ldots, B_n]$$
$$\text{then } R[A_1, \ldots, A_m] \subseteq S[B_1, \ldots, B_m]$$

- if $R[X] \subseteq S[Y]$ and $S[Y] \subseteq T[Z]$ then $R[X] \subseteq T[Z]$

Implication of inclusion dependencies

- An ID g is implied by G iff it can be derived by repeated application of the four rules shown above

- This immediately gives an expensive algorithm: apply all the rules until none are applicable.

- It turns out that the problem is inherently hard; no reasonable algorithm for solving it will ever be found

- But an important special case admits efficient solution

- Unary IDs: those of the form $R[X] \subseteq S[Y]$

- Implication $G \vdash g$ can be tested in polynomial time for unary IDs

Functional and inclusion dependencies

- In majority of applications, one has to deal with only two kinds of dependencies: FDs and IDs

- Implication problem can be solved for FDs and IDs.

- Can it be solved for FDs and IDs together?

- That is, given a set of FDs F and a set of IDs G, a FD f and an ID g. Can we determine if

$$F \cup G \vdash f$$
$$F \cup G \vdash g$$

- That is, any database that satisfies F and G, must satisfy f (or g).

- It turns out that **no** algorithm can possibly solve this problem.

Even more restrictions

- Relations are typically declared with just primary keys

- Inclusion dependencies typically occur in foreign key declarations

- Can implication be solved algorithmically for just primary keys and foreign keys?

- The answer is still NO.

- Thus, it is hard to reason about some of the most common classes of constraints found in relational databases.

When is implication solvable for FDs and IDs?

- Recall: an ID is unary if it is of the form $R[X] \subseteq S[Y]$.

- Implication can be tested for unary IDs and arbitrary FDs

- Moreover, it can be done in polynomial time.

- However, the algorithm for doing this is quite complex.

Dependencies: summary

- FDs: $X \to Y$

- Keys: $X \to U$, where U is the set of all attributes of a relation.

- Candidate key: a minimal key X; no subset of X is a key.

- Inclusion dependency: $R[A_1, \ldots, A_n] \subseteq S[B_1, \ldots, B_n]$ (unary if $n = 1$)

- Foreign key: $R[A_1, \ldots, A_n] \subseteq S[B_1, \ldots, B_n]$ and $\{B_1, \ldots, B_n\}$ is a key

- Implication problem:
    easy for FDs alone
    hard but solvable for IDs alone
    solvable by a complex algorithm for FDs and unary IDs
    unsolvable for FDs and IDs, and even
    unsolvable for primary keys and foreign keys

Database Design

- Finding database schemas with good properties

- Good properties:

    no update anomalies

    no redundancies

    no information loss

- Input: list of all attributes and constraints (usually functional dependencies)

- Output: list of relations and constraints they satisfy

Example: bad design

- Attributes: Title, Director, Theater, Address, Phone, Time, Price

- Constraints:

    FD1 Theater $\rightarrow$ Address, Phone
    FD2 Theater, Time, Title $\rightarrow$ Price
    FD3 Title $\rightarrow$ Director

- Bad design: put everything in one relation

    BAD[Title, Director, Theater, Address, Phone, Time, Price]

Why is BAD bad?

- **Redundancy**: many facts are repeated

- Director is determined by Title

    For every showing, we list both director and title

- Address is determined by Theater

    For every movie playing, we repeat the address

- **Update anomalies**:

- If Address changes in one tuple, we have inconsistency, as it must be changed in all tuples corresponding to all movies and showtimes.

- If a movie stops playing, we lose association between Title and Director

- Cannot add a movie before it starts playing

Good design

- Split BAD into 3 relations:

Relational Schema GOOD:

Table	attributes	constraints
T1	Theater, Address, Phone	FD1: Theater $\rightarrow$ Address, Phone
T2	Theater, Title, Time, Price	FD2: Theater, Time, Title $\rightarrow$ Price
T3	Title, Director	FD3: Title $\rightarrow$ Director

Why is GOOD good?

- No update anomalies:

    Every FD defines a key

- No information loss:

$$\text{T1} = \pi_{\text{Theater,Address,Phone}}(\text{BAD})$$

$$\text{T2} = \pi_{\text{Theater,Title,Time,Price}}(\text{BAD})$$

$$\text{T3} = \pi_{\text{Title,Director}}(\text{BAD})$$

$$\text{BAD} = \text{T1} \bowtie \text{T2} \bowtie \text{T3}$$

- No constraints are lost

    FD1, FD2, FD3 all appear as constraints for T1, T2, T3

Boyce-Codd Normal Form (BCNF)

- What causes updates anomalies?

- Functional dependencies $X \rightarrow Y$ where X is not a key.

- A relation is in Boyce-Codd Normal Form (BCNF) if for every nontrivial FD $X \rightarrow Y$, X is a key.

- A database is in BCNF if every relation in it is in BCNF.

Decompositions: Criteria for good design

Given a set of attributes U and a set F of functional dependencies, a *decomposition* of (U, F) is a set

$$(U_1, F_1), \ldots, (U_n, F_n)$$

where $U_i \subseteq U$ and F_i is a set of FDs over attributes U_i.

A decomposition is called:

- **BCNF** decomposition if each (U_i, F_i) is in BCNF.

Decompositions: Criteria for good design cont'd

- **lossless** if for every relation R over U that satisfies all FDs in F, each $\pi_{U_i}(R)$ satisfies F_i and

$$R \;=\; \pi_{U_1}(R) \bowtie \pi_{U_2}(R) \bowtie \ldots \bowtie \pi_{U_n}(R)$$

- **Dependency preserving** if

$$F \quad \text{and} \quad F^* = \bigcup_i F_i$$

are equivalent.

That is, for every FD f,

$$F \vdash f \quad \Leftrightarrow \quad F^* \vdash f$$

or

$$\forall X \subseteq U \quad C_F(X) \;=\; C_{F^*}(X)$$

Projecting FDs

- Let U be a set of attributes and F a set of FDs

- Let $V \subseteq U$

- Then

$$\pi_V(F) \;=\; \{X \to Y \mid X, Y \subseteq V, \quad Y \subseteq C_F(X)\}$$

- In other words, these are all FDs on V that are implied by F:

$$\pi_V(F) \;=\; \{X \to Y \mid X, Y \subseteq V, \quad F \vdash X \to Y\}$$

- Even though as defined $\pi_V(F)$ could be very large, it can often be compactly represented by a set F' of FDs on V such that:

$$\forall X, Y \subseteq V : \quad F' \vdash X \to Y \;\;\Leftrightarrow\;\; F \vdash X \to Y$$

Decomposition algorithm

Input: A set U of attributes and F of FDs
Output: A database schema $S = \{(U_1, F_1), \ldots, (U_n, F_n)\}$

1. Set $S := \{(U, F)\}$
2. While S is not in BCNF do:
   (a) Choose $(V, F') \in S$ not in BCNF
   (b) Choose nonempty disjoint X, Y, Z such that:
       (i) $X \cup Y \cup Z = V$
       (ii) $Y = C_{F'}(X) - X$
       (that is, $F' \vdash X \to Y$ and $F' \not\vdash X \to A$ for all $A \in Z$)
   (c) Replace (V, F') by

$$(X \cup Y, \pi_{X \cup Y}(F')) \;\; \text{and} \;\; (X \cup Z, \pi_{X \cup Z}(F'))$$

   (d) If there are (V', F') and (V'', F'') in S with $V' \subseteq V''$
       then remove (V', F')

Decomposition algorithm: example

Consider BAD[Title, Director, Theater, Address, Phone, Time, Price]

and constraints FD1, FD2, FD3.

Initialization: $S =$(BAD, FD1, FD2, FD3)

While loop

- Step 1: Select BAD; it is not in BCNF because
  Theater $\rightarrow$ Address, Phone
  but Theater $\not\rightarrow$ Title, Director, Time, Price
  Our sets are:
  $X = \{$Theater$\}$, $Y = \{$Address, Phone$\}$,
  $Z = \{$ Title, Director, Time, Price$\}$

Decomposition algorithm: example cont'd

- Step 1, cont'd
  $\pi_{X \cup Y}(\{$FD1,FD2,FD3$\}) = $ FD1
  $\pi_{X \cup Z}(\{$FD1,FD2,FD3$\}) = \{$FD2, FD3$\}$
  Thus, after the first step we have two schemas:
  $S_1 = (\{$Theater, Address, Phone$\}$, FD1)
  $S_1' = (\{$Theater, Title, Director, Time, Price$\}$, FD2, FD3)

- Step 2: S_1 is in BCNF, there is nothing to do.
  S_1' is not in BCNF: Title is not a key.
  Let $X = \{$Title$\}$, $Y = \{$Director$\}$, $Z = \{$Theater, Time, Price$\}$
  $\pi_{X \cup Y}(\{$FD1,FD2,FD3$\}) = $ FD3
  $\pi_{X \cup Z}(\{$FD1,FD2,FD3$\}) = $ FD2

Decomposition algorithm: example cont'd

- After two steps, we have:

  $S_1 = (\{\text{Theater, Address, Phone}\}, \text{FD1})$
  $S_2 = (\{\text{Theater, Title, Time, Price}\}, \text{FD2})$
  $S_3 = (\{\text{Title, Director}\}, \text{FD3})$

- S_1, S_2, S_3 are all in BCNF, this completes the algorithm.

Properties of the decomposition algorithm

- For any relational schema, the decomposition algorithm yields a new relational schema that is:

  in BCNF, and

  lossless

- However, the output is not guaranteed to be dependency preserving.

Surprises with BCNF

- Example: attributes Theater (th), Screen (sc), Title (tl), Snack (sn), Price (pr)

- Two FDs

$$
\begin{aligned}
\text{Theater, Screen} &\rightarrow \text{Title} \\
\text{Theater, Snack} &\rightarrow \text{Price}
\end{aligned}
$$

- Decomposition into BCNF: after one step

    (th, sc, tl; th, sc → tl),   (th, sc, sn, pr; th, sn → pr)

- After two steps:

    (th, sc, tl; th, sc → tl), (th, sn, pr; th, sn → pr), (th, sc, sn; ∅)

- The last relation looks very unnatural!

Surprises with BCNF cont'd

- The extra relation (th, sc, sn; ∅) is not needed

- Why? Because it does not add any information

- All FDs are accounted for, so are all attributes, as well as all tuples

- BCNF tells us that for any R satisfying the FDs:

$$
R \;=\; \pi_{\text{th,sc,tl}}(R) \;\bowtie\; \pi_{\text{th,sn,pr}}(R) \;\bowtie\; \pi_{\text{th,sc,sn}}(R)
$$

- However, in out case we also have

$$
R \;=\; \pi_{\text{th,sc,tl}}(R) \;\bowtie\; \pi_{\text{th,sn,pr}}(R)
$$

- This follows from the intuitive semantics of the data.

- But what is there in the schema to tell us that such an equation holds?

Multivalued dependencies

- Tell us when a relation is a join of two projections

- A multivalued dependency (MVD) is an expression of the form

$$X \longrightarrow\!\!\!\!\rightarrow Y$$

  where X and Y are sets of attributes

- Given a relation R on the set of attributes U, MVD $X \longrightarrow\!\!\!\!\rightarrow Y$ holds in it if

$$R \;=\; \pi_{XY}(R) \;\bowtie\; \pi_{X(U-Y)}(R)$$

- Simple property: $X \rightarrow Y$ implies $X \longrightarrow\!\!\!\!\rightarrow Y$

- Another definition: R satisfies $X \longrightarrow\!\!\!\!\rightarrow Y$ if for every two tuples t_1, t_2 in R with $\pi_X(t_1) = \pi_X(t_2)$, there exists a tuple t with

$$\pi_{XY}(t) \;=\; \pi_{XY}(t_1)$$
$$\pi_{X(U-Y)}(t) \;=\; \pi_{X(U-Y)}(t_2)$$

MVDs and decomposition

- If a relation satisfies a nontrivial MVD $X \longrightarrow\!\!\!\!\rightarrow Y$ and X is not a key, it should be decomposed into relations with attribute sets XY and $X(U - Y)$.

- Returning to the example, assume that we have a MVD

$$\text{Theater} \longrightarrow\!\!\!\!\rightarrow \text{Screen}$$

- Apply this to the "bad" schema (th, sc, sn; $\emptyset$) and obtain

$$(\text{th, sc}; \;\emptyset) \qquad (\text{th, sn}; \;\emptyset)$$

- Both are subsets of already produced schemas and can be eliminated.

- Thus, the FDs plus Theater $\longrightarrow\!\!\!\!\rightarrow$ Screen give rise to decomposition

$$(\text{th, sc, tl}; \;\; \text{th, sc} \rightarrow \text{tl}), \quad (\text{th, sn, pr}; \;\; \text{th, sn} \rightarrow \text{pr})$$

4th Normal Form (4NF)

- Decompositions like the one just produced are called *4th Normal Form* decompositions

- They can only be produced in the presence of MVDs

- Formally, a relation is in 4NF if for any nontrivial MVD $X \twoheadrightarrow Y$, either $X \cup Y = U$, or X is a key.

- Since $X \to Y$ implies $X \twoheadrightarrow Y$, 4NF implies BCNF

- General rule: if BCNF decomposition has "unnatural relations", try to use MVDs to decompose further into 4NF

- A useful rule: any BCNF schema that has one key that consists of a single attribute, is in 4NF

- Warning: the outcome of a decomposition algorithm may depend on the order in which constraints are considered.

BCNF and dependency preservation

- Schema Lecture[C(lass), P(rofessor), T(ime)]

- Set of FDs $F = \{$C $\to$ P, PT $\to$ C$\}$

- (Lectures, F) not in BCNF: C $\to$ P $\in F$, but C is not a key.

- Apply BCNF decomposition algorithm: $X = \{C\}$, $Y = \{P\}$, $Z = \{T\}$

- Output: $(\{C, P\}, C \to P)$, $(\{C, T\}, \emptyset)$

- We lose PT $\to$ C!

- In fact, there is no relational schema in BCNF that is equivalent to Lectures and is lossless and dependency preserving.

- Proof: there are just a few schemas on 3 attributes ... check them all ... exercise!

Third Normal Form (3NF)

- If we want a decomposition that guarantees losslessness and dependency preservation, we cannot always have BCNF

- Thus we need a slightly weaker condition

- Recall: a candidate key is a key that is not a subset of another key

- An attribute is prime if it belongs to a candidate key

- Recall (U, F) is in BCNF if for any FD $X \rightarrow A$, where $A \notin X$ is an attribute, $F \vdash X \rightarrow A$ implies that X is a key

- (U, F) is in the third normal form (3NF) if for any FD $X \rightarrow A$, where $A \notin X$ is an attribute, $F \vdash X \rightarrow A$ implies that one of the following is true:

    - either X is a key, or
    - A is prime

Third Normal Form (3NF) cont'd

- Main difference between BCNF and 3NF: in 3NF non-key FDs are OK, as long as they imply only prime attributes

- Lectures[C, P, T], $F = \{C \rightarrow P, PT \rightarrow C\}$ is in 3NF:

- PT is a candidate key, so P is a prime attribute.

- More redundancy than in BCNF: each time a class appears in a tuple, professor's name is repeated.

- We tolerate this redundancy because there is no BCNF decomposition.

Decomposition into third normal form: covers

- Decomposition algorithm needs a small set that represents all FDs in a set F

- Such a set is called *minimal cover*. Formally:

- F' is a *cover* of F iff $C_F = C_{F'}$. That is, for every f,

$$F \vdash f \quad \Leftrightarrow \quad F' \vdash f$$

- F' is a *minimal cover* if

F' is a cover,

no subset $F'' \subset F'$ is a cover of F,

each FD in F' is of the form $X \rightarrow A$, where A is an attribute,

For $X \rightarrow A \in F'$ and $X' \subset X$, $A \notin C_F(X')$.

Decomposition into third normal form: covers cont'd

- A minimal cover is a small set of FDs that give us all the same information as F.

- Example: let

$$\{A \rightarrow AB, A \rightarrow AC, A \rightarrow B, A \rightarrow C, B \rightarrow BC\}$$

- Closure:

$C_F(A) = ABC$,

$C_F(B) = BC$,

$C_F(C) = C$, and hence:

- Minimal cover: $A \rightarrow B, B \rightarrow C$

Decomposition into third normal form: algorithm

Input: A set U of attributes and F of FDs
Output: A database schema $S = \{(U_1, F_1), \ldots, (U_n, F_n)\}$

Step 1. Find a minimal cover F' of F.

Step 2. If there is $X \to A$ in F' with $XA = U$, then output (U, F').

Otherwise, select a key K, and output:

   2a) $(XA, X \to A)$ for each $X \to A$ in F', and
   2b) $(K, \emptyset)$

Decomposition into third normal form: example

- $F = \{A \to AB, A \to AC, A \to B, A \to C, B \to BC\}$

- $F' = \{A \to B, B \to C\}$

- A is a key, so output:

$$(A, \emptyset), \quad (AB, A \to B), \quad (BC, B \to C)$$

- Simplification: if one of attribute-sets produced in 2a) is a key, then step 2b) is not needed

- Hence the result is:

$$(AB, A \to B), \quad (BC, B \to C)$$

- Other potential simplifications: if we have $(XA_1, X \to A_1), \ldots,$ $(XA_k, X \to A_k)$ in the output, they can be replaced by $(XA_1 \ldots A_k, X \to A_1 \ldots A_k)$

3NF cont'd

- Properties of the decomposition algorithm: it produces a schema which is

    in 3NF

    lossless, and

    dependency preserving

- Complexity of the algorithm?

- Given the a minimal cover F', it is linear time. (Why?)

- But how hard is it to find a minimal cover?

- Naive approach: try all sets F', check if they are minimal covers.

- This is too expensive (exponential); can one do better?

- There is a polynomial-time algorithm. How?

Overview of schema design

- Choose the set of attributes U

- Choose the set of FDs F

- Find a lossless dependency preserving decomposition into:

    BCNF, if it exists

    3NF, if BCNF decomposition cannot be found

Other constraints

- SQL lets you specify a variety of other constraints:

    Local (refer to a tuple)

    global (refer to tables)

- These constraints are enforced by a DBMS, that is, they are checked when a database is modify.

- Local constraints occur in the `CREATE TABLE` statement and use the keyword `CHECK` after attribute declaration:

```
CREATE TABLE T (...
    ....
    A <type>  CHECK  <condition>,
    B <type>  CHECK  <condition,
    ....
)
```

Local constraints

- Example: the value of attribute `Rank` must be between 1 and 5:

    `Rank INT CHECK (1 <= Rank AND Rank <= 5)`

- Example: the value of attribute `A` must be less than 10, and occur as a value of attribute `C` of relation `R`:

```
A INT CHECK (A IN
    SELECT R.C FROM R WHERE R.C < 10)
```

- Example: each value of attribute `Name` occurs precisely once as a value of attribute `LastName` in relation `S`:

```
Name VARCHAR(20) CHECK (1 = SELECT COUNT(*)
                            FROM S
                            WHERE Name=S.LastName)
```

Assertions

- These assert some conditions that must be satisfied by the whole database.

- Typically assertions say that all elements in a database satisfy certain condition.

- All salaries are at least 10K:

```
CREATE ASSERTION A1 CHECK
    ( (SELECT MIN (Empl.Salary) FROM Empl >= 10000) )
```

- Most assertions use NOT EXISTS in them:
  SQL way of saying

$$\text{"every } x \text{ satisfies } F\text{"}$$

  is to say

$$\text{"does not exist } x \text{ that satisfies } \neg F\text{"}$$

Assertions cont'd

- Example: all employees of department 'sales' have salary at least 20K:

```
CREATE ASSERTION A2 CHECK
  ( NOT EXISTS (SELECT * FROM Empl
                  WHERE Dept='sales' AND Salary < 20000) )
```

- All theaters play movies that are at most 3hrs long:

```
CREATE ASSERTION A2 CHECK
  ( NOT EXISTS (SELECT * FROM Schedule S, Movies M
                  WHERE S.Title=M.Title AND M.Length > 180) )
```

Assertions cont'd

- Some assertions use counting. For example, to ensure that table T is never empty, use:

```
CREATE ASSERTION T_notempty CHECK
   ( 0 <> (SELECT COUNT(*) FROM T) )
```

- Assertions are not forever: they can be created, and dropped later:

```
DROP ASSERTION T_nonempty.
```

Triggers

- They specify a set of actions to be taken if certain event(s) took place.

- Follow the *event-condition-action* scheme.

- Less declarative and more procedural than assertions.

- Example: If an attempt is made to change the length of a movie, it should not go through.

```
CREATE TRIGGER NoLengthUpdate
AFTER UPDATE OF Length ON Movies
REFERENCING
    OLD ROW AS OldTuple
    NEW ROW AS NewTuple
FOR EACH ROW
WHEN (OldTuple.Length <> NewTuple.Length)
SET Length = OldTuple.Length
WHERE title = NewTuple.title
```

Analysis of the trigger

- `AFTER UPDATE OF Length ON Movies` specifies the **event**: Relation `Movies` was modified, and attribute `Length` changed its value.

- `REFERENCING`
      `OLD ROW AS OldTuple`
      `NEW ROW AS NewTuple`

  says how we refer to tuples before and after the update.

- `FOR EACH ROW` – the tigger is executed once for each update row.

- `WHEN (OldTuple.Length <> NewTuple.Length)` – **condition** for the trigger to be executed (Lenght changed its value).

- `SET Length = OldTuple.Length`
  `WHERE title = NewTuple.title`

  is the **action**: Lenght must be restored to its previous value.

Another trigger example

- Table `Empl(emp_id,rank,salary)`

- Requirement: the average salary of managers should never go below $100,000.

- Problem: suppose there is a complicated update operation that affects many tuples. It may initially decrease the average, and then increase it again. Thus, we want the trigger to run after all the statements of the update operation have been executed.

- Hence, `FOR EACH ROW` trigger cannot work here.

Another trigger example cont'd

Trigger for the previous example:

```
CREATE TRIGGER MaintainAvgSal
AFTER UPDATE OF Salary ON Empl
REFERENCING
    OLD TABLE AS OldTuples
    NEW TABLE AS NewTuples
FOR EACH STATEMENT
WHEN ( 100000 > (SELECT AVG(Salary)
                FROM Empl WHERE Rank='manager') )
BEGIN
    DELETE FROM Empl
    WHERE (emp_id, rank, salary) in NewTuples;
    INSERT INTO Empl
        (SELECT * FROM OldTuples)
END;
```

Analysis of the trigger

- `AFTER UPDATE OF Salary ON Empl` specifies the **event**: Relation `Empl` was modified, and attribute `Salary` changed its value.

- `REFERENCING`
    ```
    OLD TABLE AS OldTuples
    NEW TABLE AS NewTuples
    ```
  says that we refer to the set of tuples that were inserted into `Empl` as `NewTuples` and to the set of tuples that were deleted from `Empl` as `OldTuples`.

- `FOR EACH STATEMENT` – the tigger is executed once for the entire update operation, not once for each updated row.

- `WHEN ( 100000 > (SELECT AVG(Salary)`
        `FROM Empl WHERE Rank='manager') )`
  is the **condition** for the trigger to be executed: after the entire update, the average `Salary` of managers is less than $100K.

Analysis of the trigger cont'd

- `BEGIN`
  ```
      DELETE FROM Empl
      WHERE (emp_id, rank, salary) in NewTuples;
      INSERT INTO Empl
            (SELECT * FROM OldTuples)
  END;
  ```
  is the **action**, that consists of two updates between `BEGIN` and `END`.

-
  ```
      DELETE FROM Empl
      WHERE (emp_id, rank, salary) in NewTuples;
  ```
  deletes all the tuples that were inserted in the illegal update, and

-
  ```
      INSERT INTO Empl
            (SELECT * FROM OldTuples)
  ```
 re-inserts all the tuples that were deleted in that update.

Conceptual modeling

- How does one come up with a set of attributes and FDs?

- Often by using conceptual design techniques, typically represented by entity-relationship models

- Database describes objects (entities), relationships between them, and their attributes.

- Goal: describe entities, relationships, and attributes based on a real-world situation that needs to be modeled by a database, find relational database design, and then apply normalization techniques.

- Steps:

  describe entities

  describe relationships between them
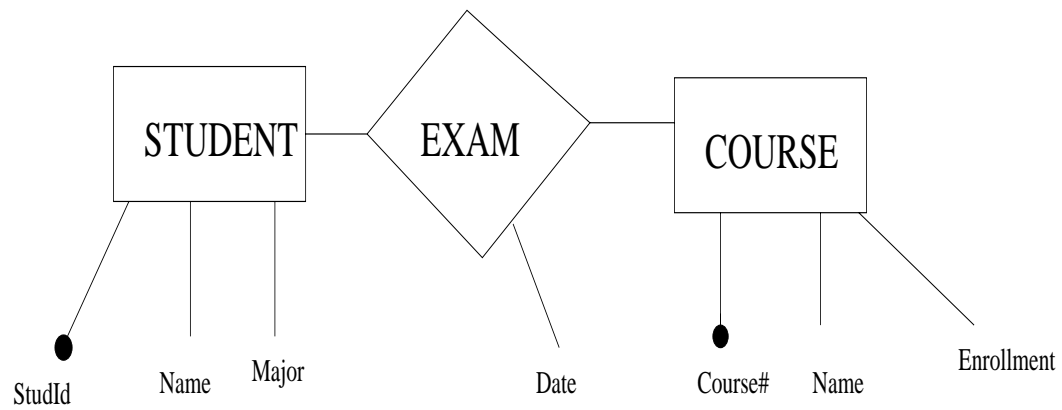
  describe attributes

  translate into relational model

ENTITIES

STUDENT		COURSE	

EMPLOYEE		PROJECT	

RELATIONSHIPS

STUDENT — EXAM — COURSE

EMPLOYEE — WORKS — PROJECT

ATTRIBUTES

Translating into relational model: rules

- Each entity E_i is described by a relation T_i

- Key attributes of E_i form a key of T_i

- Each relationship R_j is described by a relation S_j

- For a relationship R_j between $E_1, \ldots, E_k$, the key of S_j is the set of all attributes which are keys for $T_1, \ldots, T_k$

- Sometimes attributes have to be renamed

TRANSLATION INTO RELATIONS

Student(StudId, Name, Major)   Course(Course#, Name, Enrollment)

Exam(StudId,Course#,Date)

Database design in SQL

```
CREATE TABLE Student
       (StudId char(10) not null primary key,
        Name char(20), Major char(5))

CREATE TABLE Course
       (Course# int not null primary key,
        Name char(30), Enrollment int)

CREATE TABLE Exam
       (Student char(10) not null,
        Course int not null,
        Date int,
        unique (Student, Course))
```